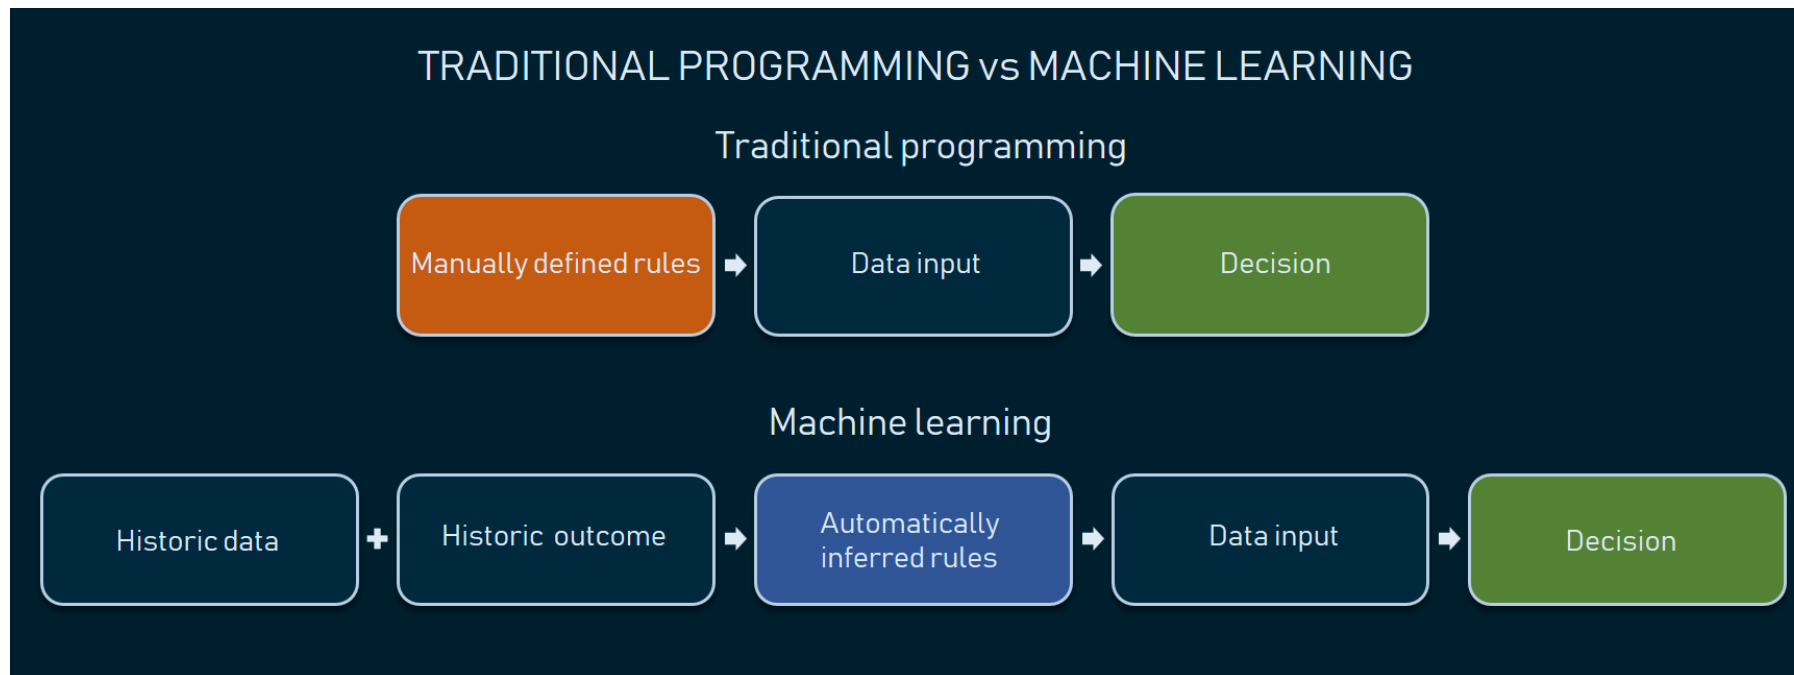


# Neural Network in One Hour

Ravi Shankar

# Traditional Programming




# Traditional Programming vs Machine Learning

- ▶ If  $(a > b)$  return  $a$ ; else  $b$
- ▶ Look at the inputs and outputs


a	b	output
10	5	10
4	8	8
4	4	4
6	2	?



## Traditional Programming

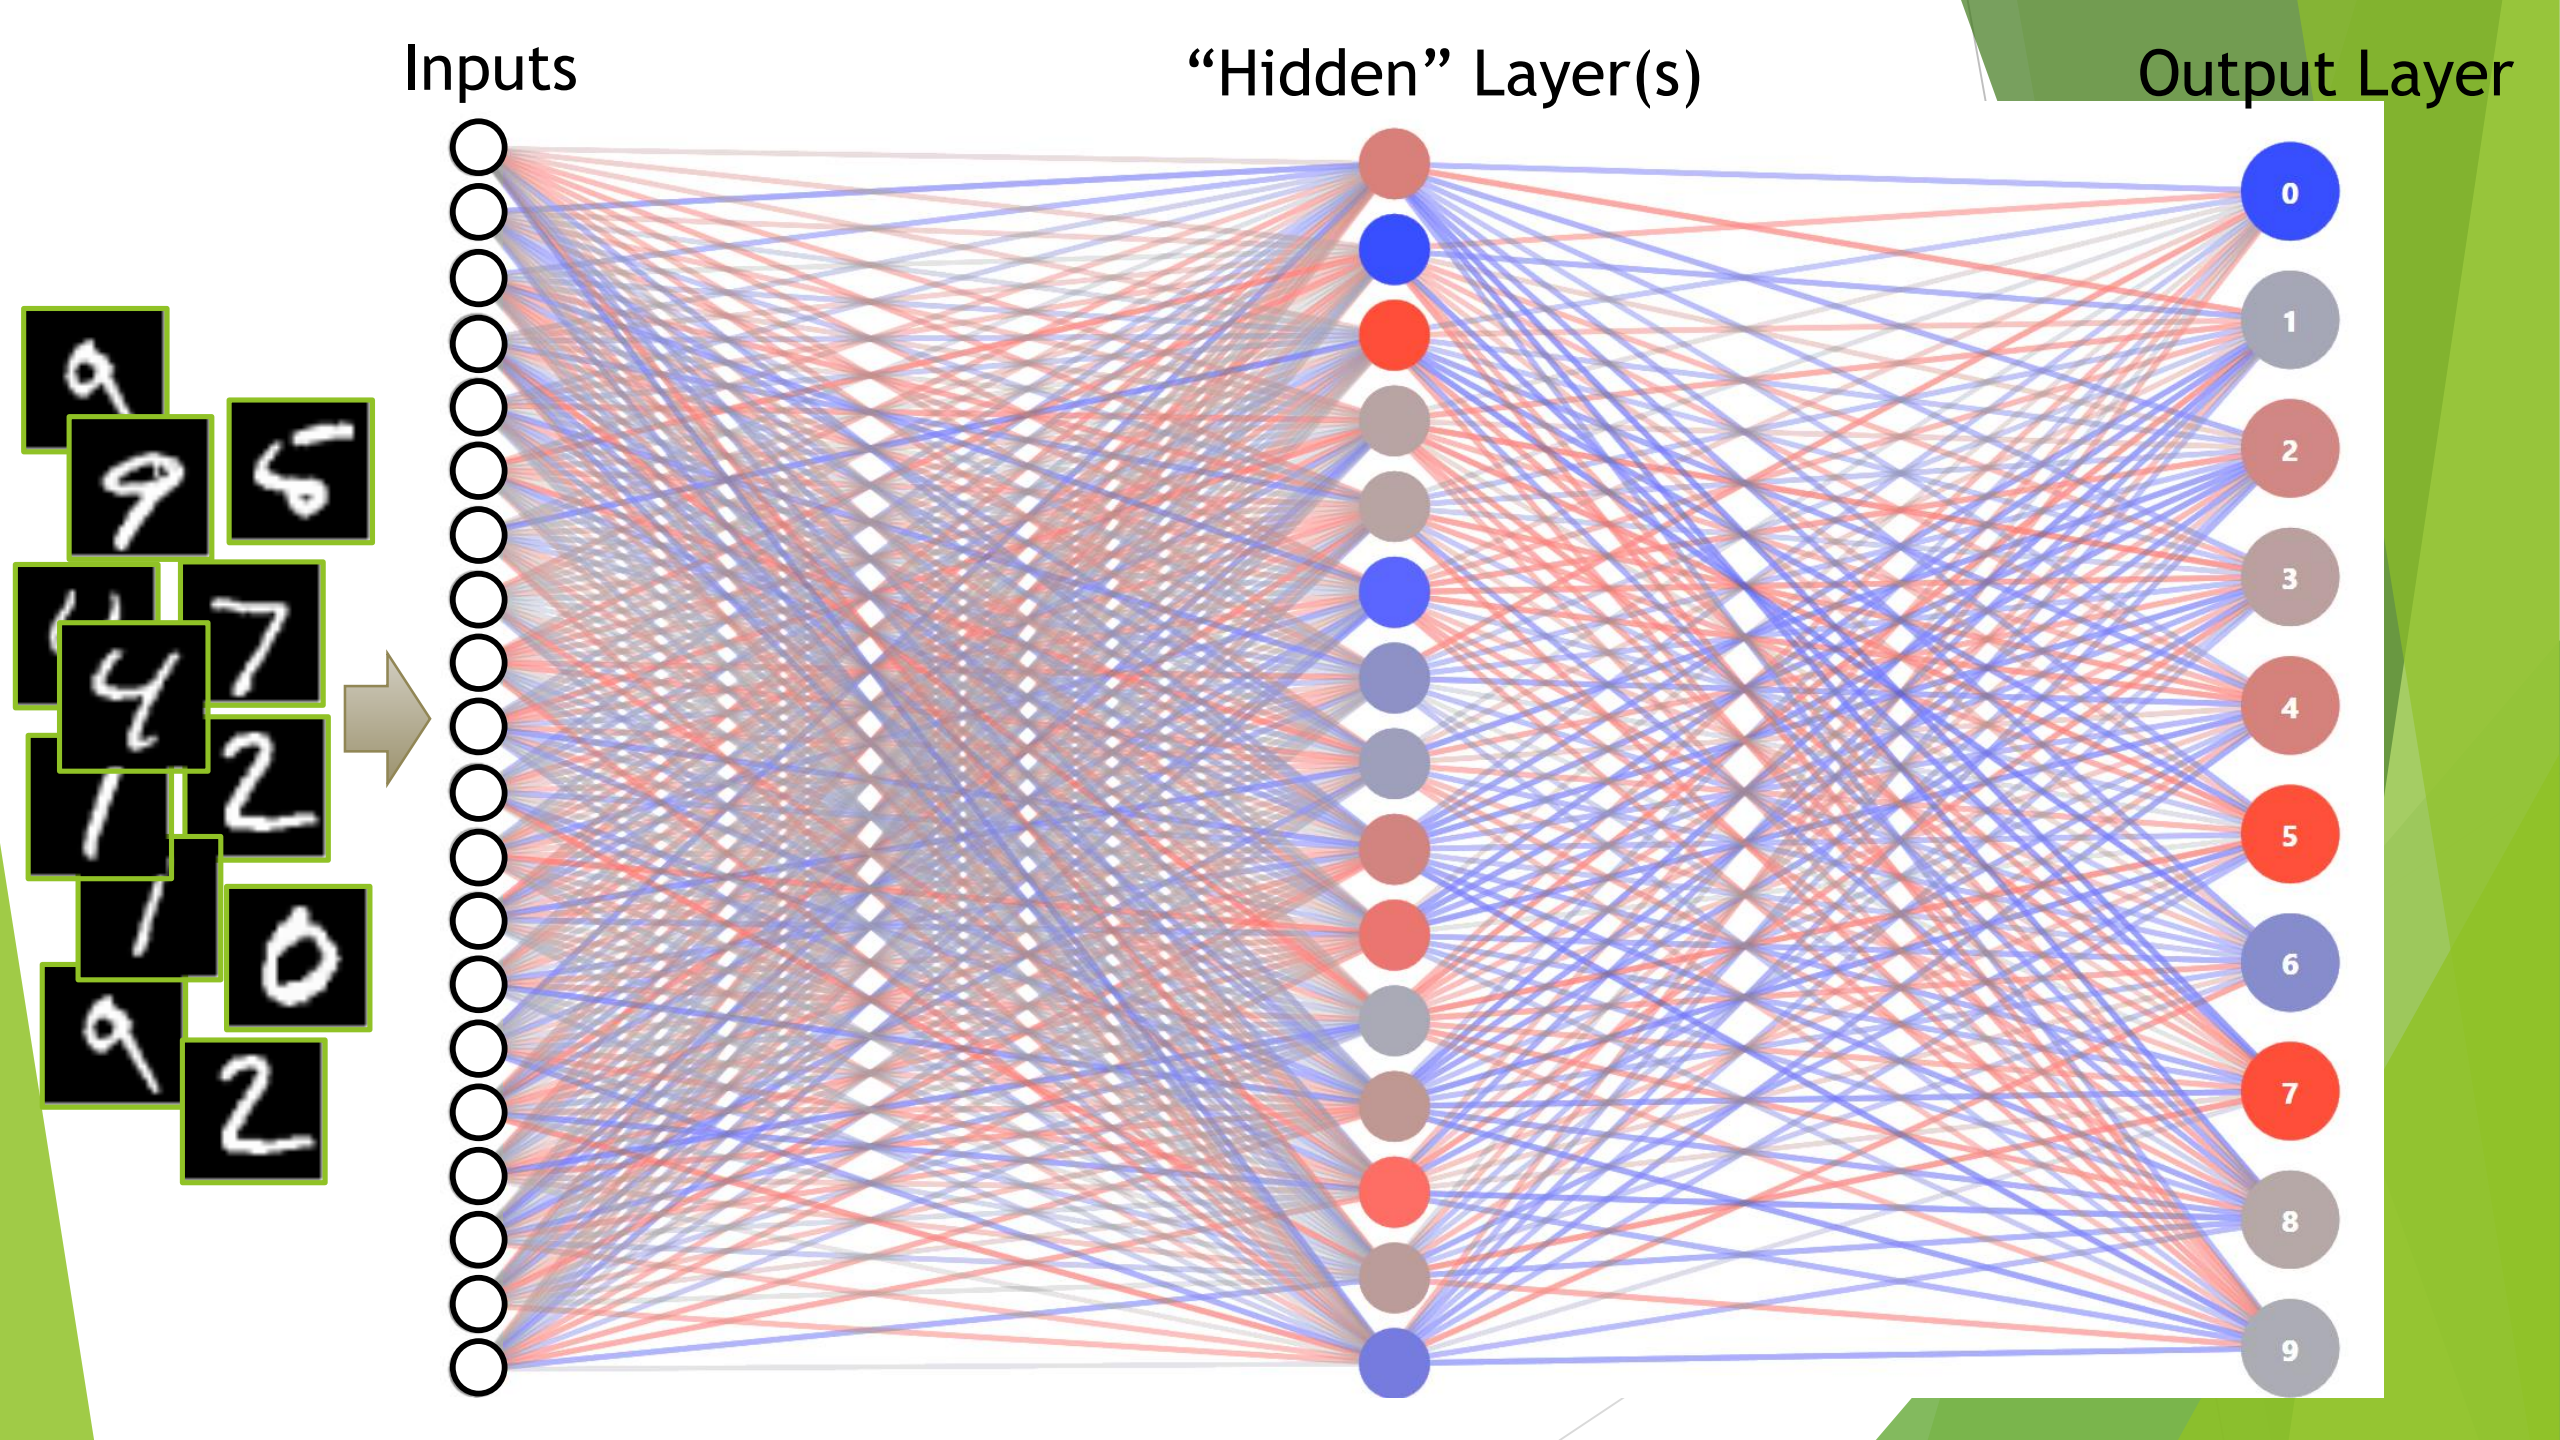


## Machine Learning



**MODEL !!**









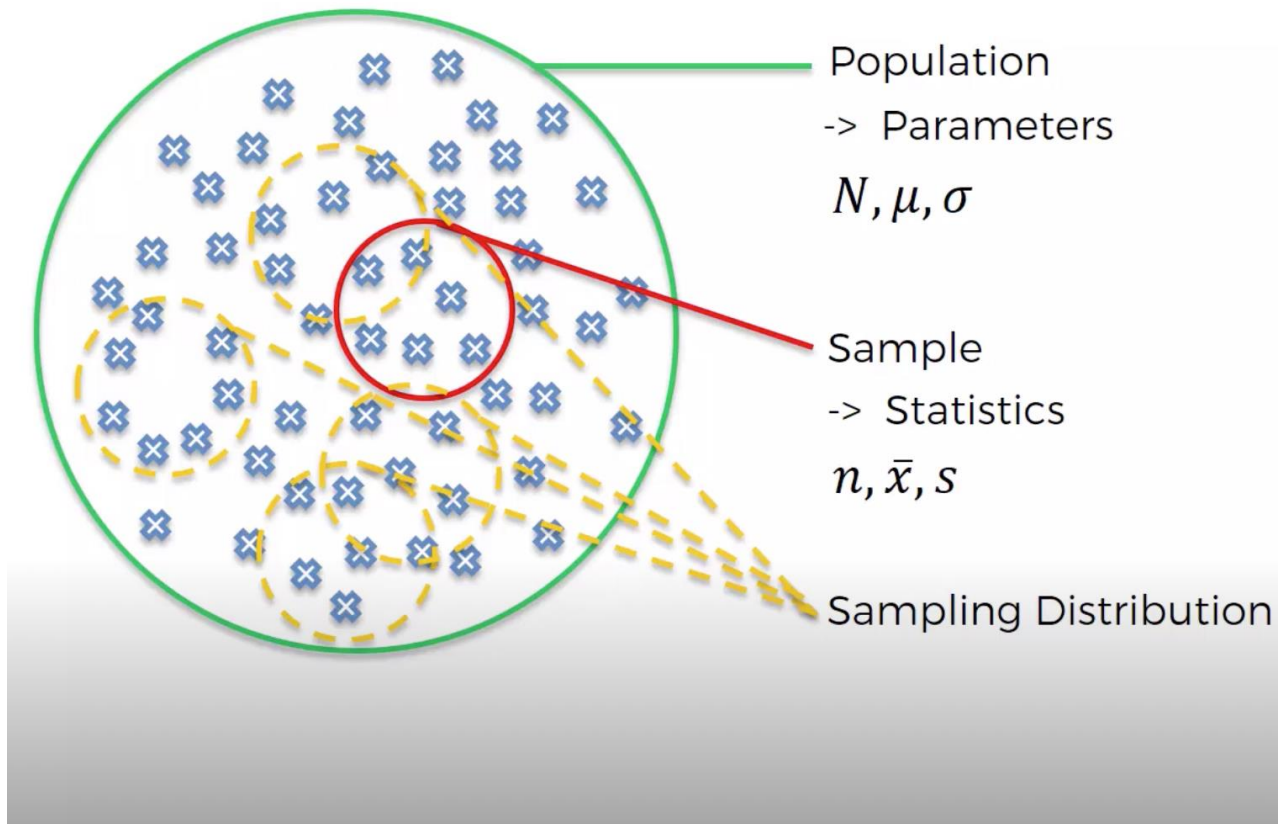


# Epoch

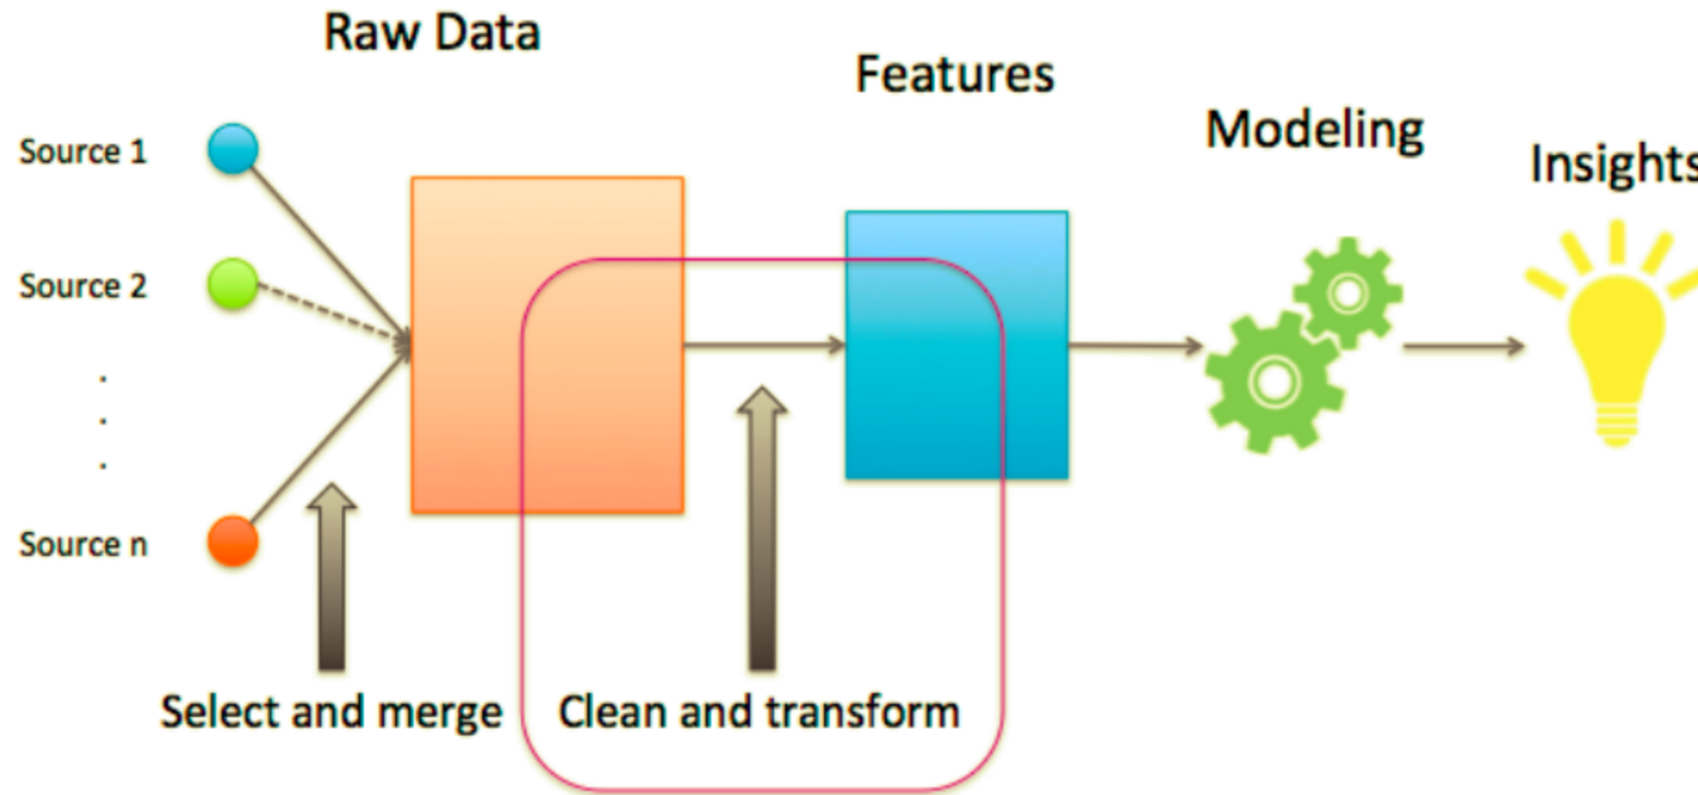
The number of epochs is a hyperparameter that defines the number times that the learning algorithm will work through the entire training dataset.



# Population and Samples

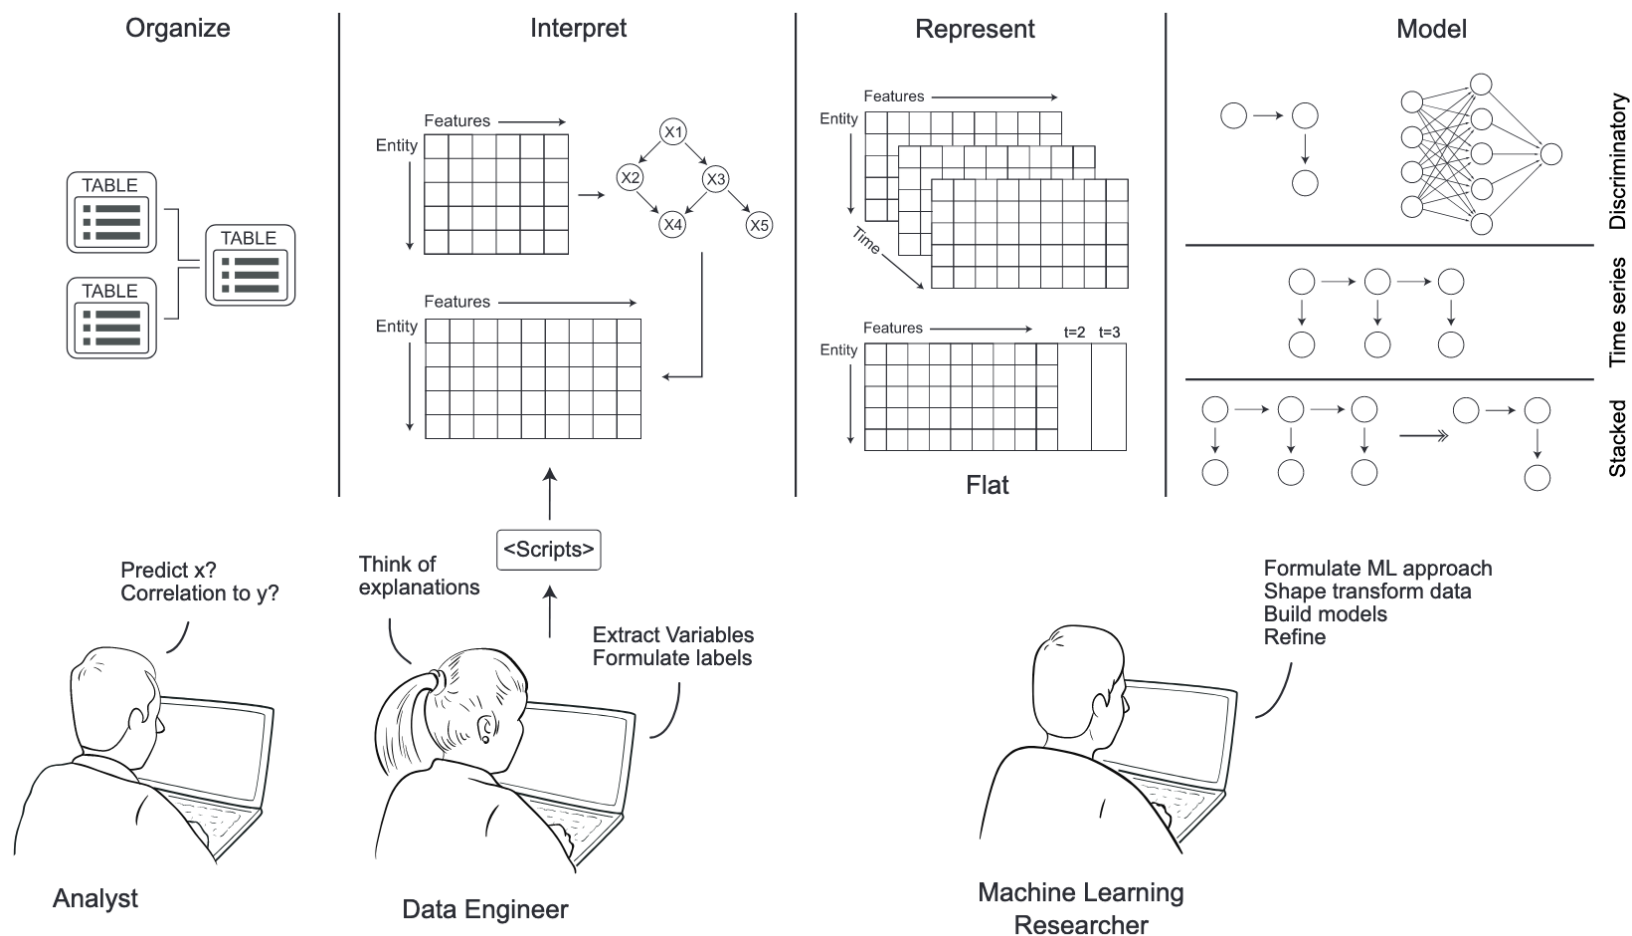


# Ideal Input Identification



# Deep Feature Synthesis (DFS)

You can automate data science endeavors - AutoML plays less role, but auto featuring is amazing





# Deep Feature Synthesis (DFS)

Deep Feature Synthesis algorithm, which is capable of generating features that express a rich feature-space. We can develop an end-to-end Data Science Machine which can:


- (a) automatically generate features via Deep Feature Synthesis
- (b) autotune a machine learning pathway to extract the most value out of the synthesized features
- (c) produce submissions for online data science competitions. Matched human level performance when competing in data science competitions using the Data Science Machine

[http://www.jmaxkanter.com/static/papers/DSAA\\_DSM\\_2015.pdf](http://www.jmaxkanter.com/static/papers/DSAA_DSM_2015.pdf)

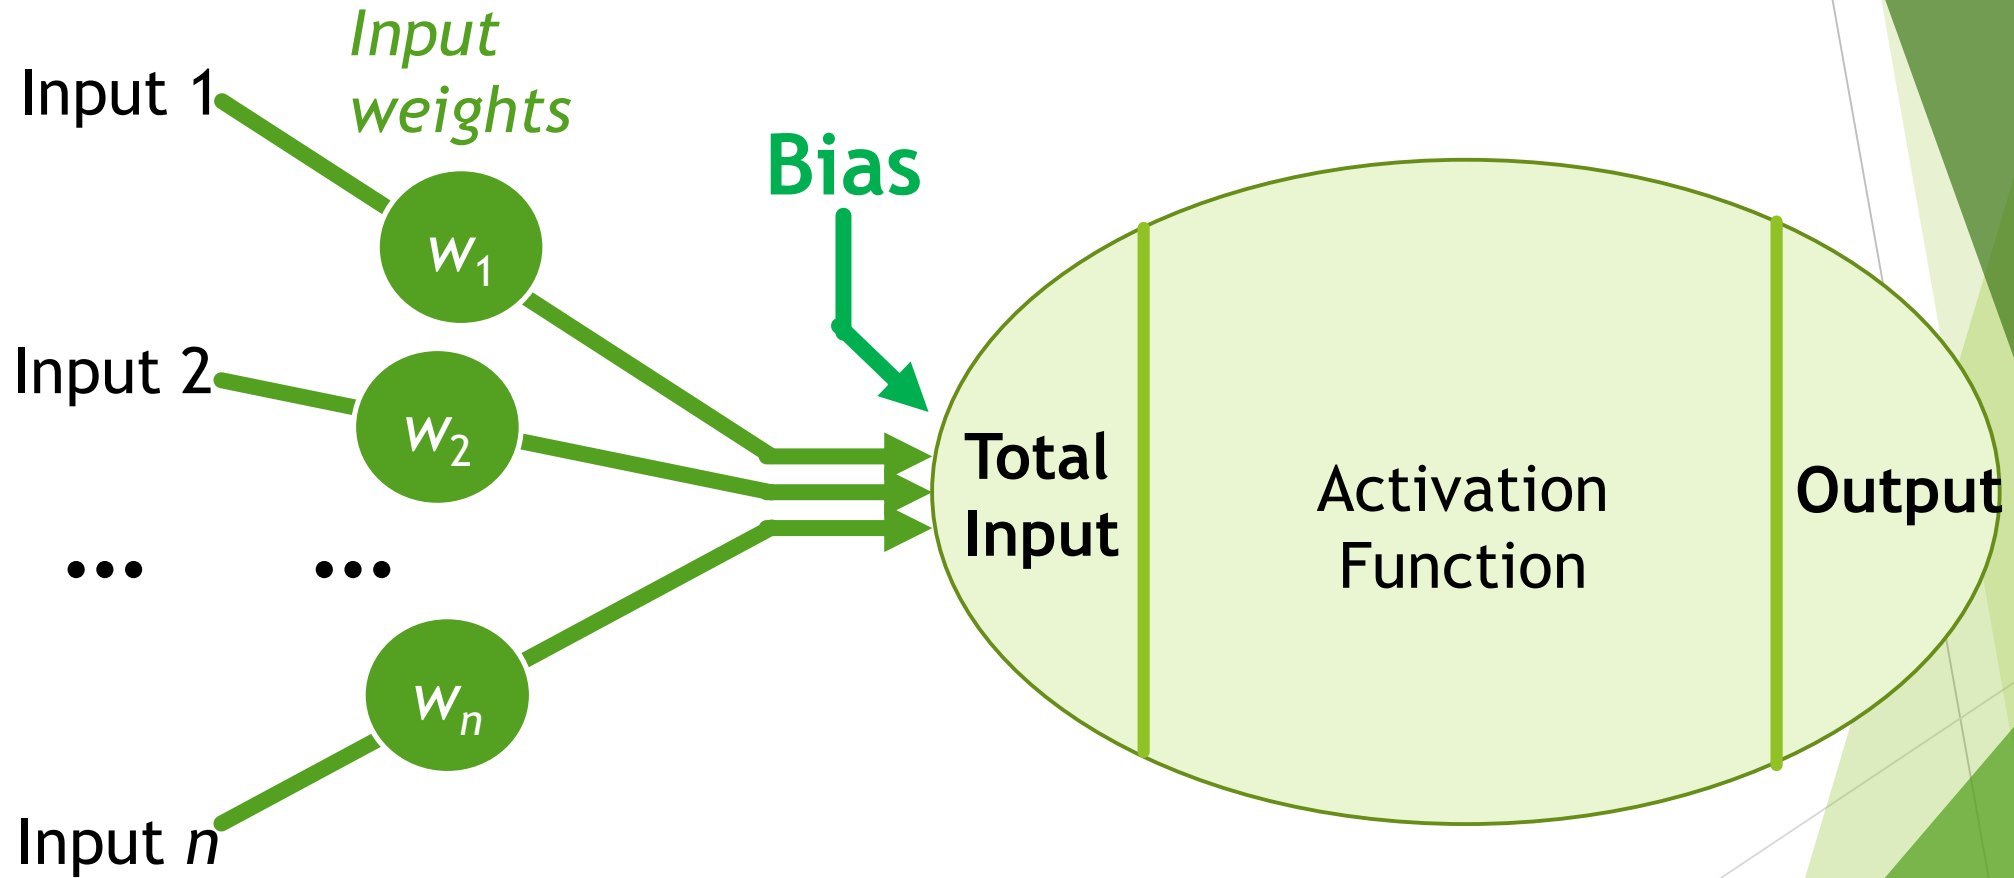
# The Best Places to End Up

- ▶ Many options available
  - ▶ TensorFlow
  - ▶ Microsoft Cognitive Toolkit
  - ▶ Keras
  - ▶ Pytorch
  - ▶ Caffe
  - ▶ Anaconda

# Inspired by Human Brain




# The AI Neuron




All values are numbers

# Comparator






Better



Logistic Sigmoid

$$y = \frac{1}{1 + e^{-x}}$$

More Better




ReLU

*Rectified Linear Unit*

```
double Relu (double x) =>  
  x >= 0 ? x : 0;
```

# Leaky ReLU

## *Rectified Linear Unit*



```
double Relu (double x) =>  
    x >= 0 ? x : x / 100;
```

# Some Adjustments

Rather than:

$$\text{Loss} = \text{sum}(\text{error})^2$$

We *can* consider:

$$\text{Loss} = \textcolor{red}{1/2} \text{ sum}(\text{error})^2$$

- ▶ We could instead use average (mean)
- ▶ *Cost* is a synonym for loss

**We don't usually calculate the loss!**

# Stochastic Gradient Descent

foreach training sample:

Determine slope of the loss function at  
current position


(the error tells us the slope!)

Step the following amount:

- slope \* learning rate



# Architecture



Ideal



## Future: GANs

# Apple hires Google AI expert Ian Goodfellow to direct machine learning

Goodfellow is best known for inventing generative adversarial networks (GANs), which pair two AI algorithms together with the goal of continuously improving one another — one AI could be tasked with creating realistic images, while the other AI acts as a judge of real versus fake images, such that both AIs spur each other to become better over time. His research also involved combating adversarial attacks on neural networks that could undermine a GAN's ability to perform its functions, including methods that were undetectable to human observers.

Thank You & CodeOntheBeach-2019